



Probability Calibration

Output of classifiers

- The class
- The probability of being of class 1



Probability and confidence

When performing classification we may want to predict the class label, and also to obtain a probability, certainty or confidence around the respective label.



Probability as certainty


- Probabilities can be much more informative than labels:
- “The model predicts this claim is fraudulent” vs “The model predicts this claim is 90% likely to be fraudulent”
- To convey likelihood, we need calibrated probabilities

Calibrated Probability: what is it?

- In calibrated probabilities, the probability reflects the true likelihood
- Calibration is the concordance of predicted probabilities with the occurrence of positive cases
- If 10 observations obtain a probability of 0.8 and the probability is calibrated, I expect around 8 of those to belong to the positive class

Calibration Curve


- If the probability is calibrated, we should see a match between the number of positive cases and the predicted probability



<https://amueller.github.io/COMS4995-s20/slides/aml-10-calibration-imbalanced-data/#5>

Calibration Curve

- If the probability is calibrated, we should see a match between the number of positive cases and the predicted probability



<https://amueller.github.io/COMS4995-s20/slides/aml-10-calibration-imbalanced-data/#5>


ML models and probability

- Logistic regression returns calibrated probabilities by design
- Some machine learning models return uncalibrated probabilities
 - Decision trees
 - Naïve Bayes
- Some models do not support probability prediction

ML models and probability

- Logistic regression returns calibrated probabilities by design
- Boosted algorithms, Random Forests and Naïve Bayes, push probability mass away from 0 and 1 yielding sigmoid shape distortion in the predicted probabilities
- Some models output the class with no probability, eg, SVMs

ML models and probability



<https://scikit-learn.org/stable/modules/calibration.html>

Uncalibrated probability: should we care?

To determine model performance, not necessarily.

As users of the model, we might.

- ✓ This patient is 80% likely to have cancer is more useful to a doctor than 0 or 1, the class prediction.
- ✓ This claim is 70% likely to be fraudulent is more useful to fraud investigators than 0.2, the uncalibrated prediction.

THANK YOU

www.trainindata.com